

**Ex no: 8**

## **Hadoop - Map Reduce**

### **Setup a Single Node Hadoop cluster and show all the process through WEB UI.**

**Aim:**

To setup a Single Node Hadoop cluster and show all the process through WEB UI.

**Procedure:**

1. Install java version 8 or higher.
2. In the environment variables create a variable in the user var
3. iables and name it as “JAVA\_HOME” and specify the JDK file path in the variable value.
4. To check whether Java’s path is set properly open command prompt and type “java -version”.
5. Java’s version will be displayed if the java configuration is done properly.
6. Now to install Hadoop goto the official Apache Hadoop website and download the latest version.
7. Extract the tar file using WinRAR into the C drive.
8. After extracting the Hadoop file, open the environment variables and create a new variable “HADOOP\_HOME” in the user variable with the value of the variable as the file path of the Hadoop’s bin folder
9. In the Hadoop file create a new folder called “data” inside which you should create two more folders one is “namenode” and the other is “datanode” folders.
10. To configure Hadoop edit the following files which are located in the etc folder :
  - core-site.xml
  - hdfs-site.xml
  - mapred-site.xml
  - yarn-site.xml
11. After editing the above files open hadoop-env file in notepad and set the java home location to it.
12. Now goto command prompt and type the command “hdfs namenode -format”. This command is used to format the HDFS (Hadoop Distributed File System) namenode. This command initializes the directory structure of the namenode by creating the necessary file system metadata.
13. NameNode works on the Master System. The primary purpose of Namenode is to manage all the MetaData. Metadata is the list of files stored in HDFS(Hadoop Distributed File System). As we know the data is stored in the form of blocks in a Hadoop cluster. So the DataNode on which or the location at which that block of the file is stored is mentioned in MetaData. All information regarding the logs of

the transactions happening in a Hadoop cluster (when or who read/wrote the data) will be stored in MetaData. MetaData is stored in the memory.

14. Now type the commands “start-dfs.cmd” and “start-yarn.cmd”. This command will start all the daemons at once.
15. Daemons mean Process. Hadoop Daemons are a set of processes that run on Hadoop.
16. Goto the browser and type “localhost:9870” this will open Hadoop in the browser.

## Output:

Edit User Variable

Variable name:

JAVA\_HOME

Variable value:

C:\Program Files\Java\jdk-22

Browse Directory...

Browse File...

OK

Cancel

```
Microsoft Windows [Version 10.0.22631.3958]
(c) Microsoft Corporation. All rights reserved.

C:\Users\mannu>java -version
java version "22.0.1" 2024-04-16
Java(TM) SE Runtime Environment (build 22.0.1+8-16)
Java HotSpot(TM) 64-Bit Server VM (build 22.0.1+8-16, mixed mode, sharing)

C:\Users\mannu>
```

## Download

Hadoop is released as source code tarballs with corresponding binary tarballs for convenience. The downloads are distributed via mirror sites and should be checked for tampering using GPG or SHA-512.

Version	Release date	Source download	Binary download	Release notes
3.4.0	2024 Mar 17	<a href="#">source (checksum signature)</a>	<a href="#">binary (checksum signature)</a> <a href="#">binary-aarch64 (checksum signature)</a>	<a href="#">Announcement</a>
3.3.6	2023 Jun 23	<a href="#">source (checksum signature)</a>	<a href="#">binary (checksum signature)</a> <a href="#">binary-aarch64 (checksum signature)</a>	<a href="#">Announcement</a>
2.10.2	2022 May 31	<a href="#">source (checksum signature)</a>	<a href="#">binary (checksum signature)</a>	<a href="#">Announcement</a>

To verify Apache Hadoop® releases using GPG:

1. Download the release `hadoop-X.Y.Z-src.tar.gz` from a [mirror site](#).
2. Download the signature file `hadoop-X.Y.Z-src.tar.gz.asc` from [Apache](#).
3. Download the [Hadoop KEYS](#) file.
4. `gpg --import KEYS`
5. `gpg --verify hadoop-X.Y.Z-src.tar.gz.asc`



<div> <div> <div></div> <div>This PC</div> </div> <div> <div></div> <div>Windows (C:)</div> </div> <div> <div></div> <div>hadoop-3.4.0</div> </div> <div> <div></div> <div>data</div> </div> </div>			
<div> <div> <div></div> <div></div> <div></div> <div></div> <div></div> </div> <div>Sort</div> <div>View</div> <div></div> </div>			
Name	Date modified	Type	Size
<div> <div></div> <div>datanode</div> </div>	05-08-2024 07:59	File folder	
<div> <div></div> <div>namenode</div> </div>	04-08-2024 15:16	File folder	

```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

## core-site.xml

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.data.dir</name>
<value>file:///C:/hadoop-3.4.0/data/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>file:///C:/hadoop-3.4.0/data/datanode</value>
</property>
</configuration>
```

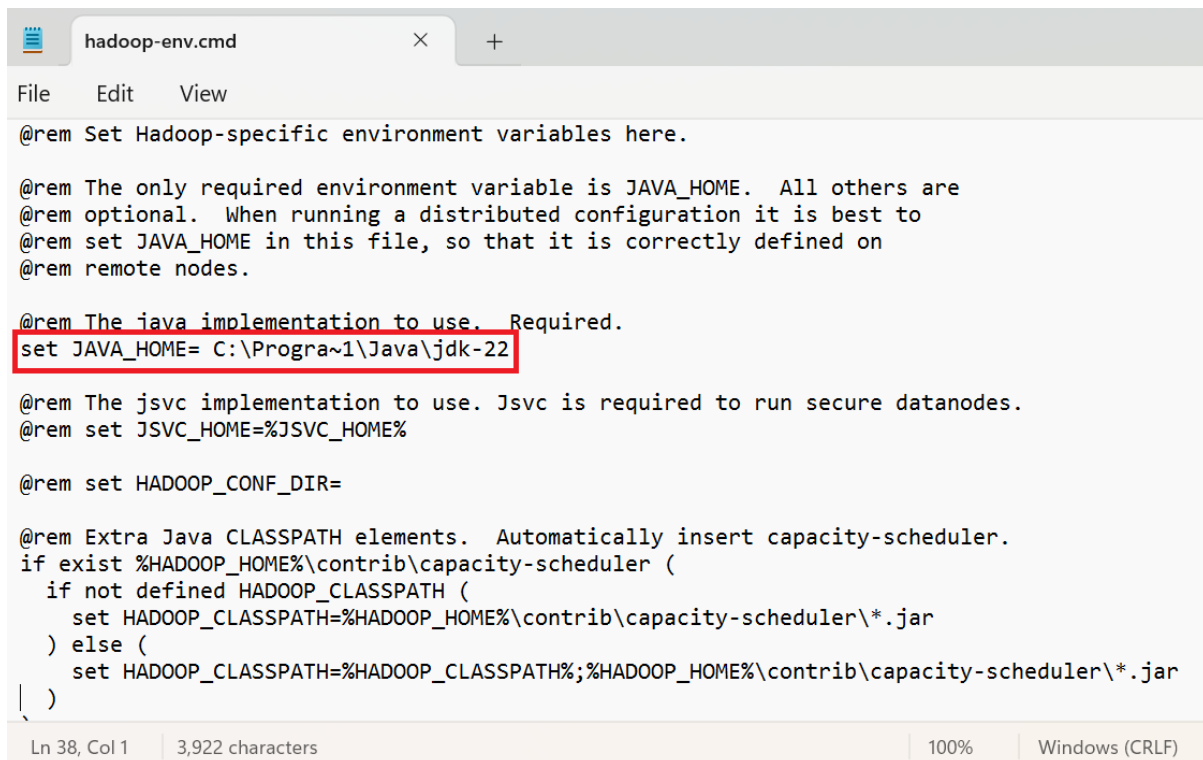
## hdfs-site.xml

```
<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
```

## mapred-site.xml

```
<configuration>
<!-- Site specific YARN configuration properties -->
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
</configuration>
```

## yarn-site.xml



```
hadoop-env.cmd
File Edit View
@rem Set Hadoop-specific environment variables here.

@rem The only required environment variable is JAVA_HOME. All others are
@rem optional. When running a distributed configuration it is best to
@rem set JAVA_HOME in this file, so that it is correctly defined on
@rem remote nodes.

@rem The java implementation to use. Required.
set JAVA_HOME= C:\Progra~1\Java\jdk-22

@rem The jsvc implementation to use. Jsvc is required to run secure datanodes.
@rem set JSVC_HOME=%JSVC_HOME%

@rem set HADOOP_CONF_DIR=

@rem Extra Java CLASSPATH elements. Automatically insert capacity-scheduler.
if exist %HADOOP_HOME%\contrib\capacity-scheduler (
  if not defined HADOOP_CLASSPATH (
    set HADOOP_CLASSPATH=%HADOOP_HOME%\contrib\capacity-scheduler\*.jar
  ) else (
    set HADOOP_CLASSPATH=%HADOOP_CLASSPATH%;%HADOOP_HOME%\contrib\capacity-scheduler\*.jar
  )
)
```

Ln 38, Col 1 | 3,922 characters | 100% | Windows (CRLF)

## hadoop-env file

```
Command Prompt
Microsoft Windows [Version 10.0.22631.3958]
(c) Microsoft Corporation. All rights reserved.

C:\Users\mannu>hdfs namenode -format
2024-08-07 23:06:44,240 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = Shreeya/192.168.1.10
STARTUP_MSG: args = [-format]
STARTUP_MSG: version = 3.4.0
STARTUP_MSG: classpath = C:\hadoop-3.4.0\etc\hadoop;C:\hadoop-3.4.0\share\hadoop\common;C:\hadoop-3.4.0\share\hadoop\com
mon\lib\animal-sniffer-annotations-1.17.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\audience-annotations-0.12.0.jar;C:\
hadoop-3.4.0\share\hadoop\common\lib\avro-1.9.2.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\bcprov-jdk15on-1.70.jar;C:\
hadoop-3.4.0\share\hadoop\common\lib\checker-qual-2.5.2.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-beanutils-1.
9.4.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-cli-1.5.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-co
dec-1.15.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-collections-3.2.2.jar;C:\hadoop-3.4.0\share\hadoop\common\l
ib\commons-compress-1.24.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-configuration2-2.8.0.jar;C:\hadoop-3.4.0\
share\hadoop\common\lib\commons-daemon-1.0.13.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-io-2.14.0.jar;C:\hado
op-3.4.0\share\hadoop\common\lib\commons-lang3-3.12.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-logging-1.2.jar
;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-math3-3.6.1.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-net-3.9
.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\commons-text-1.10.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\curator-c
lient-5.2.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\curator-framework-5.2.0.jar;C:\hadoop-3.4.0\share\hadoop\common\
lib\curator-recipes-5.2.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\dnsjava-3.4.0.jar;C:\hadoop-3.4.0\share\hadoop\com
mon\lib\failureaccess-1.0.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\gson-2.9.0.jar;C:\hadoop-3.4.0\share\hadoop\common
\lib\guava-27.0-jre.jar;C:\hadoop-3.4.0\share\hadoop\common\lib\hadoop-annotations-3.4.0.jar;C:\hadoop-3.4.0\share\hado
```

```
Command Prompt
Microsoft Windows [Version 10.0.22631.3958]
(c) Microsoft Corporation. All rights reserved.

C:\Users\mannu>start-dfs.cmd

C:\Users\mannu>start-yarn.cmd
starting yarn daemons

C:\Users\mannu>
```

Namenode information

localhost:9870/dfshealth.html#tab-overview

Hadoop

Overview

Datanodes

Datanode Volume Failures

Snapshot

Startup Progress

Utilities

## Overview 'localhost:9000' (✓active)

Started:	Thu Aug 08 07:55:07 +0530 2024
Version:	3.4.0, rbd8b77f398f626bb7791783192ee7a5dfeec760
Compiled:	Mon Mar 04 12:05:00 +0530 2024 by root from (HEAD detached at release-3.4.0-RC3)
Cluster ID:	CID-3c7fc0a5-fa2c-4b2e-b157-b1a3214a892f
Block Pool ID:	BP-275814143-192.168.1.10-1723052211560

## Summary

Security is off.

Safemode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).

Heap Memory used 62.7 MB of 80 MB Heap Memory. Max Heap Memory is 1000 MB.  
Non Heap Memory used 52.45 MB of 54.88 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	352.98 GB
Configured Remote Capacity:	0 B
DFS Used:	321 B (0%)
Non DFS Used:	154.21 GB
DFS Remaining:	198.77 GB (56.31%)
Block Pool Used:	321 B (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion (including replicas)	0

Block Deletion Start Time	Thu Aug 08 07:55:07 +0530 2024
Last Checkpoint Time	Thu Aug 08 07:53:35 +0530 2024
Last HA Transition Time	Never
Enabled Erasure Coding Policies	RS-6-3-1024k

## NameNode Journal Status

Current transaction ID: 2

Journal Manager	State
FileJournalManager(root=\tmp\hadoop-mannu\dfs\name)	EditLogFileOutputStream(\tmp\hadoop-mannu\dfs\name\current\edits_inprogress_0000000000000000002)

## NameNode Storage

Storage Directory	Type	State
\tmp\hadoop-mannu\dfs\name	IMAGE_AND_EDITS	Active

## DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	352.98 GB	321 B (0%)	198.77 GB (56.31%)	321 B	1

Hadoop, 2024.

### Result:

Thus to setup a Single Node Hadoop cluster and show all the process through WEB UI has been executed successfully.